

TOOLKIT 1 OF 4

AI Procurement Questions for Hiring Teams

48 structured vendor evaluation questions across 6 categories, with red flags, guidance notes, and a comparison framework.

Introduction

How to use this toolkit

This toolkit is for HR Directors, Talent Acquisition leads, Procurement teams, and Ops or Legal stakeholders evaluating AI hiring tools. It is written for people who are not AI experts. The goal is to give non-technical buyers the language and structure to ask better questions and recognise weak answers.

Four steps

Step 1. Work through Section 2 (Internal alignment) with your team before any vendor conversation. These prompts surface assumptions and gaps that are better resolved before you are in a demo.

Step 2. Select the categories most relevant to your context. All six matter, but weight them based on your industry, candidate volume, and regulatory exposure.

Step 3. Use the questions live in vendor calls, not just as pre-sent documents. How a vendor responds in real time tells you as much as the answer itself.

Step 4. Use the comparison sheet after each vendor conversation to build a side-by-side record of your evaluation.

Core and deeper-dive questions

Each category labels questions as Core or Deeper Dive.

CORE	Prioritise in every vendor conversation. Ask these first.
DEEPER DIVE	For follow-up calls, or when Legal or Risk are in the room.

What this toolkit is not

This is not a contract template, a legal opinion, or a substitute for legal review. It is a structured vendor evaluation tool. For contract negotiation, data processing agreements, and formal compliance review, work with qualified legal counsel.

A vendor who cannot answer these questions clearly is answering you anyway.

Section 2

Internal alignment before vendor conversations

These prompts are for your team, not for vendors. Work through them before your first vendor call. They are not counted in the 48 vendor questions. They exist to make sure your organisation arrives at procurement conversations with a clear internal position.

- 1 Do we have a named person accountable for AI use in our hiring process?
- 2 Do we understand which parts of our current process already involve AI, including tools embedded in our ATS?
- 3 Have we defined what we are trying to achieve with this tool, beyond saving time?
- 4 Do we know which protected characteristics are most relevant to our candidate population and hiring context?
- 5 Have we discussed with Legal or Compliance what our obligations are under GDPR and AI regulation for this use case?
- 6 Do we have a process for candidates to raise concerns about how AI is used in our hiring process?
- 7 Are our hiring managers and recruiters prepared to explain AI use to candidates if asked?
- 8 Have we agreed internally on what meaningful human oversight looks like in our specific process?

NOTE

If your team cannot agree on the answers to these prompts, that is important information. It means you are not yet ready to evaluate vendors. Resolve internal alignment first.

Category 1 of 6

Bias and fairness

Why this category matters

AI hiring tools trained on historical data can encode and scale existing biases in ways that are hard to detect after the fact. Most vendors will tell you they have tested for bias. The important question is how, with what methodology, across which protected characteristics, and what the results actually showed. This category helps you move past reassurance and into evidence.

Questions to ask

- CORE** What protected characteristics does your bias testing cover, and how frequently is testing conducted?
- CORE** Can you share the results of your most recent bias audit, including any disparate impact findings?
- CORE** What methodology do you use to test for bias? Is it internal testing, third-party audit, or both?
- CORE** If bias is identified in the outputs, what is your process for addressing it and notifying clients?
- CORE** Has your tool ever been found to produce discriminatory outcomes in practice? If so, what happened?
- DEEPER DIVE** How does the model perform across candidates with non-traditional backgrounds or career gaps?
- DEEPER DIVE** What data was the model originally trained on, and how representative is it of diverse candidate populations?
- DEEPER DIVE** Do clients receive ongoing fairness reporting, and how does performance vary across different role families?

WHAT A STRONG ANSWER LOOKS LIKE

The vendor provides specific protected characteristics tested, shares a summary of results, names the methodology, and can describe what remediation looks like when issues are found.

RED FLAGS AND WEAK ANSWERS

Vague assurances such as 'we take bias very seriously' without supporting evidence. Reluctance to share audit results. Claims that the tool is 'bias-free' should raise concern rather than reassurance.

Category 2 of 6

Data and privacy

Why this category matters

AI hiring tools process significant amounts of personal data about candidates, often including inferred attributes candidates are unaware of. Understanding what data is collected, how long it is retained, whether it trains future models, and who has access is not just a compliance question. It is a candidate trust question. This category surfaces practices most buyers do not ask about until something goes wrong.

Questions to ask

- CORE** What personal data does your tool collect from candidates, and is any of it inferred rather than directly provided?
- CORE** Is candidate data from our hiring process used to train or improve your models? Can we opt out?
- CORE** How long is candidate data retained, and what is the deletion process when requested?
- CORE** Which third parties or subprocessors have access to candidate data, including any external model providers?
- CORE** How do you handle data subject access requests from candidates under GDPR or equivalent regulation?
- DEEPER DIVE** Who within your organisation has access to candidate data from our account?
- DEEPER DIVE** Where is candidate data stored geographically, and does it transfer outside the UK or EU?
- DEEPER DIVE** What happens to our candidate data if we end our contract with you?

WHAT A STRONG ANSWER LOOKS LIKE

Clear, specific answers. The vendor can point to their data processing agreement, has a documented deletion process, allows opt-out of training data use, and can name their subprocessors and data residency arrangements.

RED FLAGS AND WEAK ANSWERS

Inability to confirm whether data is used for model training. Vague references to industry-standard security. No clear answer on data residency, deletion timelines, or which subprocessors have access.

Category 3 of 6

Explainability and auditability

Why this category matters

If a candidate asks why they were not progressed, can you give them a meaningful answer? If a decision is challenged, can you reconstruct how it was reached? Many AI hiring tools produce outputs without a clear explanation of the inputs or logic involved. Explainability is both a data protection consideration and a basic standard of fairness. This category helps your legal and procurement teams assess the risks before you sign.

Questions to ask

- CORE** How does your tool score or rank candidates, and what factors carry the most weight?
- CORE** Which parts of the scoring process can you explain clearly, and which are less interpretable?
- CORE** Can you explain a specific hiring decision in plain language if asked by a candidate or internal reviewer?
- CORE** What documentation is produced for each candidate assessment, and how long is it retained?
- DEEPER DIVE** What logs, decision records, or audit trails can we access directly as a client, rather than only on request?
- DEEPER DIVE** How have you designed this tool to manage the risks around automated decision-making under data protection law?
- DEEPER DIVE** If a candidate requests an explanation of a decision, what can you provide and in what timeframe?
- DEEPER DIVE** Has your tool's decision-making process ever been subject to external audit? What did that involve?

WHAT A STRONG ANSWER LOOKS LIKE

The vendor can describe in plain terms what the model scores and why. They have documentation practices in place, can support candidate explanation requests, give clients direct access to audit trails, and have considered their position on automated decision-making under data protection law.

RED FLAGS AND WEAK ANSWERS

Deflection on how the model works, citing proprietary IP. No process for candidate explanation requests. Uncertainty about data protection obligations. Claims that the model is just a tool and all decisions are human, while also describing automated shortlisting.

Category 4 of 6

Human oversight and escalation

Why this category matters

Meaningful human oversight is not a checkbox that a recruiter clicks without reviewing the reasoning behind an automated output. This category establishes whether human review is genuinely embedded in the process, whether candidates have a route to escalate, whether recruiters can override the tool, and whether your organisation has defined accountability for what the system does in practice.

Questions to ask

- | | | |
|---|-------------|---|
| 1 | CORE | At which specific points in the hiring process does human review take place? |
| 2 | CORE | Can a candidate request human review of an automated assessment? If so, how? |
| 3 | CORE | Can human reviewers override the tool's output, and is that override recorded for audit purposes? |
| 4 | CORE | Who in our organisation is accountable for decisions made with the assistance of this tool? |
| 5 | CORE | What safeguards exist to prevent hiring teams from rubber-stamping automated outputs? |
| 6 | DEEPER DIVE | If the tool's output contradicts a recruiter's assessment, what guidance do you provide? |
| 7 | DEEPER DIVE | What training do you provide or recommend for hiring teams using this tool? |
| 8 | DEEPER DIVE | How do you ensure human review points are genuinely substantive rather than procedural? |

WHAT A STRONG ANSWER LOOKS LIKE

The vendor can identify specific named points where human review occurs, has a candidate escalation process, confirms overrides are tracked, provides recruiter training, and is clear that the tool supports rather than replaces human decision-making.

RED FLAGS AND WEAK ANSWERS

Describing 'human in the loop' without specifying where or how. No candidate escalation process. Framing the tool as autonomous in early stages. Inability to confirm whether overrides are tracked or what recruiter training looks like.

Category 5 of 6

Vendor accountability

Why this category matters

When an AI hiring tool produces a bad outcome, who is responsible? Vendors often design contracts to limit their own liability while placing accountability on the buyer. This category helps you understand what the vendor will actually stand behind, what recourse exists when something goes wrong, and whether accountability is genuinely shared or quietly transferred to your organisation at the point of signing.

Questions to ask

- CORE** How do you approach liability and client support if the tool contributes to a harmful or discriminatory outcome?
- CORE** What is your incident response process if a bias, security, or model-performance issue is discovered?
- CORE** How quickly would we be notified if an issue affecting our hiring process were identified?
- CORE** What ongoing support do you provide to help clients use the tool responsibly and compliantly?
- DEEPER DIVE** Do you carry professional indemnity or product liability insurance for AI-influenced hiring decisions?
- DEEPER DIVE** Have any clients faced legal challenges related to the use of your tool? How were they handled?
- DEEPER DIVE** What contractual protections exist if a tool failure leads to a decision that causes harm?
- DEEPER DIVE** How do you track changes in employment law and AI regulation, and how do you communicate these to clients?

WHAT A STRONG ANSWER LOOKS LIKE

Clear liability position, documented incident response with defined notification timelines, relevant insurance, and ongoing compliance support. They can reference how real issues have been handled, not just hypothetical processes.

RED FLAGS AND WEAK ANSWERS

Immediate redirection to the legal team when liability is mentioned. Statements that all decisions are the client's responsibility. No incident response process. No clear notification timeline when issues are discovered.

Category 6 of 6

Contract and exit

Why this category matters

Buying an AI hiring tool is easier than leaving one. Data becomes embedded, processes are built around outputs, and exit terms are rarely scrutinised at signing. This category ensures your organisation understands what happens to your data if you leave, what portability looks like in practice, what transition support is available, and whether you are entering a relationship that is genuinely reversible.

Questions to ask

- 1 **CORE** What happens to all candidate data held by your platform when we end the contract?
- 2 **CORE** Can we export our historical data in a usable format before we leave? What does that look like?
- 3 **CORE** What is the minimum contract term, and what are the notice requirements or penalties for early exit?
- 4 **CORE** What support do you provide during offboarding or migration to another vendor?
- 5 **DEEPER DIVE** Do you use our data or hiring outcomes to improve your models after we leave?
- 6 **DEEPER DIVE** Are there technical or commercial arrangements that would make switching to a different vendor difficult?
- 7 **DEEPER DIVE** What does the full off-boarding process look like, and how long does complete data deletion take?
- 8 **DEEPER DIVE** If your company is acquired or ceases to operate, what protections exist for our data and continuity?

WHAT A STRONG ANSWER LOOKS LIKE

A documented exit process with defined timelines for data deletion and export. Data in portable formats. No punitive exit terms. A clear answer on post-contract data use. Transition support included or available.

RED FLAGS AND WEAK ANSWERS

Vague deletion timelines. No exportable format for historical data. Long minimum terms with significant exit penalties. No transition support. Uncertainty about whether your data continues to train their models after you leave.

Section 9

Vendor comparison sheet

Complete this sheet after each vendor conversation. Score each category 1 to 5, where 1 is poor or evasive and 5 is clear, specific, and evidenced.

ON WEIGHTING

Not all categories need equal weight. Adjust scoring emphasis to reflect your context. High-volume hiring may weight bias most. Regulated industries may weight explainability and contract terms most. Decide this before you compare vendors.

Category	Vendor A	Vendor B	Vendor C	Notes
Bias and fairness				
Data and privacy				
Explainability and auditability				
Human oversight and escalation				
Vendor accountability				
Contract and exit				
Overall score (out of 30)				
Recommended for next stage?				

Scoring guide

5 = Clear, specific, evidenced. No follow-up needed. 4 = Good, minor gaps. 3 = Partial. Follow-up required. 2 = Vague or inconsistent. Significant concern. 1 = Evasive, unable to answer, or raised a red flag.

Section 10

Final notes

This toolkit is designed to make vendor conversations more structured, more productive, and more honest. It is not designed to make procurement slower or to turn every vendor call into an interrogation.

Vendors with well-built tools tend to welcome rigorous questions. They know their product holds up. Evasiveness in response to reasonable questions is itself useful information.

Responsible procurement starts with better questions. The teams that ask them are already differentiated.

What this toolkit is and is not

This is a structured vendor evaluation tool. It is not legal advice, a contract template, or a guarantee of compliance. The questions are prompts for procurement and legal review, not statements of legal obligation. For contract negotiation and data processing agreements, work with qualified legal counsel.

The questions reflect current best practice in responsible AI procurement as of April 2026. Regulatory requirements are evolving. Review this toolkit annually and update your procurement process as obligations change.

Part of the AI Hiring Toolkit

This document is one of four practical frameworks for organisations navigating AI in hiring responsibly. The full toolkit also includes the AI Hiring Policy Framework, Bias Audit Checklist, and Candidate Transparency Guide.

aievolvingyou.com/resources