

AI Hiring Policy Framework

Assess your current governance, identify gaps, and build a draft AI Hiring Policy ready for internal review.

INCLUDES

- 42 governance review questions
- Internal alignment prompts
- AI tool inventory template
- Draft AI Hiring Policy template

AI Hiring Policy Framework

What this toolkit is for This toolkit helps HR Directors, Heads of Talent, and People Operations leaders build a responsible AI hiring policy. It is the final toolkit in the AI Hiring Toolkit series and is designed to be the practical conclusion of the governance work the series supports.

Who it is for HR and People Operations teams. DPOs and Data Protection leads. General Counsel and Employment Law teams. Anyone responsible for deploying or governing AI tools in a hiring process.

What it produces Part 1 is a structured policy readiness review across six governance categories. Part 2 is a draft AI Hiring Policy template that organisations can adapt, circulate internally, and take to Legal or the board for review. The objective is to leave the buyer with a usable first draft of their AI Hiring Policy, not simply an assessment of their current state.

The EU AI Act's high-risk AI obligations, including those covering recruitment and selection tools, become fully applicable on 2 December 2027. This was confirmed in June 2026, delayed from the original 2 August 2026 date. A separate, general AI transparency rule still applies from 2 August 2026. The UK's Data (Use and Access) Act 2025 mandatory complaints procedure came into force on 19 June 2026. Organisations without a documented AI hiring policy are exposed to regulatory, legal, and reputational risk that is no longer theoretical.

Procurement. Transparency. Audit. Policy. Each toolkit in the series produces something your organisation can use immediately. This one produces the governance document that ties everything together.

AI Hiring Policy Framework

How to use this toolkit

This toolkit has two parts. Part 1 is a policy readiness review: six categories of questions that help you identify the gaps in your current governance before you write or update your policy. Part 2 is a draft AI Hiring Policy template that you can complete, adapt, and take to Legal or the board for review.

Work through Part 1 with your governance team first. Use what you find to inform how you complete Part 2. The template includes guidance notes and sample text for each section. Replace the sample text with language that reflects your organisation, your tools, and your actual process.

Step 1	Work through the internal alignment prompts on the next page before starting Part 1.
Step 2	Complete Part 1 with your HR, Legal, DPO, and DEI leads. Note every gap honestly.
Step 3	Use the gaps identified in Part 1 to prioritise which sections of the policy template to complete first.
Step 4	Complete Part 2 using the sample text as a starting point. Adapt every section to your organisation.
Step 5	Circulate the draft policy for internal review. Take it to Legal, your DPO, and executive sign-off.

Question types

CORE	Minimum standard questions. Every CORE gap represents meaningful legal or governance exposure and should be addressed before the policy is signed off.
DEEPER DIVE	Extended governance questions. Important for organisations scaling AI use, operating in multiple jurisdictions, or preparing for external scrutiny.

Estimated completion time	Part 1: 90 to 120 minutes with your governance team. Part 2: 2 to 4 hours to complete the policy template, depending on your current documentation.
Primary owner	CHRO / Head of Talent / HR Director
Suggested contributors	DPO or Data Protection Lead, General Counsel or Employment Law, CISO or IT Security, DEI Lead, Head of Recruitment Operations, Procurement
Suggested review cycle	Full review annually, and immediately following any significant regulatory change, AI tool change, or candidate complaint.

What this toolkit does not do

This toolkit does not provide legal advice. It does not replace specialist counsel on the EU AI Act, UK GDPR, or employment discrimination law. The policy template is a starting point, not a finished legal document. Have it reviewed by your legal team before final approval.

Before you start

Internal alignment prompts

Do not skip this section. These prompts surface the gaps that matter most before you begin.

Work through these prompts with your HR, Legal, and DPO leads before starting the six categories. They are not counted in the question total.

Can we name every AI tool currently used in our hiring process, including features built into our ATS?

If the answer involves any uncertainty, building the tool inventory in Category 2 is your first priority. You cannot govern what you cannot name.

Who in our organisation would you call if a candidate complained tomorrow that an AI tool had rejected them unfairly?

If the answer is unclear or involves multiple people with no clear lead, governance ownership in Category 1 is a gap that needs closing before anything else.

Have we completed a Data Protection Impact Assessment (DPIA) for every AI hiring tool we currently use?

A DPIA is a structured assessment used to identify and mitigate privacy and data protection risks before deploying a high-risk processing activity such as AI-assisted candidate screening. DPIAs are legally required, not optional. The ICO has found them absent or inadequate in the majority of organisations it has audited. If the answer is no or uncertain, Category 4 is a priority.

If a regulator asked us to evidence our AI hiring governance tomorrow, what would we show them?

Documents, named owners, review dates, audit trails, and training records. If the answer is a verbal description of what we generally do, this toolkit is for you.

Do our vendor contracts address what happens when the model is updated, and who is responsible if it produces a discriminatory outcome?

Most do not. Category 5 covers the contractual safeguards that turn governance intentions into enforceable commitments.

Have we told our recruiters and hiring managers what meaningful human review actually requires them to do?

Knowing a tool is in use is not the same as knowing how to review its output critically. Category 3 covers what that training and documentation needs to include.

Governance ownership and decision rights

Why this category matters

The ICO's 2024 recruitment AI audit consistently identified absent or unclear internal ownership as a structural governance failure. Organisations that cannot name who is responsible for their AI hiring tools cannot govern them. CIPD (January 2026) positions HR as the natural owner of AI hiring governance, as an extension of existing employment relationship responsibilities. The EU AI Act Article 26 goes further: deployers of high-risk AI must assign oversight to named individuals with the authority and skills to act.

Questions to review

1	CORE	Is there a named role accountable for our AI hiring policy and its implementation?
2	CORE	Does that person have the authority to approve new AI tools, pause existing ones, and escalate concerns?
3	CORE	Have we identified a named DPO or data protection lead who co-owns governance for AI in hiring?
4	CORE	Do we have a cross-functional review group covering HR, Legal, DPO, IT Security, and DEI for AI hiring decisions?
5	DEEPER DIVE	Is there a documented decision rights matrix showing who approves, reviews, and can override AI tool use?
6	DEEPER DIVE	Do board or executive leaders have visibility of AI hiring tools and the risks they carry?

WHAT GOOD LOOKS LIKE

- A named HR Director or CHRO is documented as policy owner, with the authority to approve, pause, or remove AI tools.
- A cross-functional group meets at least twice a year to review AI hiring tools, incidents, and regulatory changes.
- Decision rights are documented: who approves new tools, who signs off DPIAs, who can halt a tool.

RED FLAGS AND WEAK PRACTICE

- Nobody can clearly name who owns AI hiring governance in practice.
- Governance is split across HR, Legal, and IT with no single accountable person.
- AI tools have been deployed without any formal approval or sign-off process.
- The board or executive team has no visibility of AI hiring tools or their risks.

AI tool inventory and classification

Why this category matters

Governance cannot be applied to tools that are not documented. The DSIT Responsible AI in Recruitment guidance identifies building an AI tool inventory as the essential first step in any governance programme. The EU AI Act requires deployers to classify each AI system against Annex III: tools that score, rank, or profile candidates are almost certainly high-risk, triggering a significantly higher level of obligation. Tools used informally by individual recruiters outside any approved process represent a shadow AI risk that most organisations have not mapped.

Questions to review

1	CORE	Do we have a documented inventory of every AI tool currently used in our hiring process?
2	CORE	Does the inventory include ATS algorithmic features, not just standalone AI assessment platforms?
3	CORE	Have we classified each tool as high-risk or lower-risk under the EU AI Act framework?
4	CORE	Is there an approval process that any new AI hiring tool must go through before use?
5	DEEPER DIVE	Have we identified informal or ad-hoc AI use by individual recruiters, such as using public generative AI tools with candidate data?
6	DEEPER DIVE	Is the inventory reviewed and updated at least annually, and when tools change materially?

WHAT GOOD LOOKS LIKE

- A documented inventory exists covering all AI tools in use, including ATS features and any generative AI tools used in the hiring process.
- Each tool has been classified against EU AI Act Annex III criteria, with the classification rationale documented.
- A formal approval process exists for new tools, with DPO and Legal sign-off required before deployment.

RED FLAGS AND WEAK PRACTICE

- No inventory exists. Tools have been deployed without central tracking.
- The ATS's built-in ranking or scoring features have not been included in any governance review.
- Individual recruiters are using public AI tools with candidate data and this is not covered by policy.
- Classification decisions have not been documented or reviewed.

Human oversight standards

Why this category matters

The ICO's April 2026 fairness spotlight and its 2024 outcomes report both identify rubber-stamping as the most common and consequential human oversight failure. A recruiter who views an AI score and accepts it without scrutiny may be making a decision that is unlawfully solely automated under UK GDPR and the DUAA 2025. The EU AI Act requires deployers of high-risk AI to assign oversight to people with the necessary skills, authority, and support, including the ability to halt the system. The policy must translate these regulatory concepts into specific, trainable steps.

Questions to review

1	CORE	Have we defined what meaningful human review means in our process, beyond a recruiter clicking approve?
2	CORE	Does our human review process require the reviewer to read the underlying application or assessment before confirming a shortlist decision?
3	CORE	Can human reviewers override AI output, and is that override recorded in the system?
4	CORE	Are reviewers trained on what the AI tool is measuring and how to interpret its output critically?
5	CORE	Is there a process for candidates to request human review of an AI-influenced decision?
6	DEEPER DIVE	Do we have controls to prevent time pressure or volume pressure from reducing the quality of human review?
7	DEEPER DIVE	Is the human review process documented clearly enough to be auditable if challenged?

WHAT GOOD LOOKS LIKE

- A written definition of meaningful human review is in the policy, specifying the minimum steps a reviewer must take.
- Overrides are tracked and reviewers document their rationale when they depart from an AI recommendation.
- Reviewers have completed training on what the tool measures and how to challenge its output.
- Candidates have a clear, accessible route to request human review.

RED FLAGS AND WEAK PRACTICE

- Human review exists procedurally but is not defined. Reviewers receive the AI score with no requirement to examine the underlying material.
- Overrides are not tracked. There is no record of whether human review is substantive.
- Recruiters have not been trained on what the tool measures or how to interpret its output.
- Candidates are not told that human review is available or how to request it.

Data protection obligations

Why this category matters

The ICO's 2024 recruitment AI audit found widespread data protection failures: excessive data collection, indefinite retention, purpose-limitation breaches, and unlawful use of candidate data to train third-party AI models. DPIAs are legally required where processing is likely to result in high risk. The ICO is clear that AI used for automated screening meets this threshold. The Data (Use and Access) Act 2025, partially in force from February 2026, updated the UK ADM framework. A mandatory complaints procedure came into force on 19 June 2026.

Questions to review

1	CORE	Has a DPIA been completed for every AI hiring tool before deployment?
2	CORE	Do we have a documented lawful basis for each category of candidate data processed by AI tools?
3	CORE	Are retention schedules in place for AI outputs and candidate scoring data, with no indefinite retention?
4	CORE	Have we confirmed that candidate data is not being used by vendors to train their central AI models without our knowledge?
5	CORE	Is our internal complaints procedure for AI-influenced decisions documented and operational?
6	DEEPER DIVE	For EU operations: have we completed or planned a Fundamental Rights Impact Assessment for high-risk AI tools?
7	DEEPER DIVE	Are DPIAs reviewed and updated when AI tools change materially or vendors update their models?

WHAT GOOD LOOKS LIKE

- DPIAs have been completed for all AI hiring tools, with DPO sign-off, before deployment.
- Retention schedules are specific, documented, and enforced. No AI output data is held indefinitely.
- Vendor contracts explicitly prohibit use of candidate data for model training without a separate lawful basis.
- A documented complaints procedure is operational and candidates are informed of their rights.

RED FLAGS AND WEAK PRACTICE

- DPIAs have not been completed, or were completed retrospectively and not reviewed since.
- Retention is managed by default vendor settings with no specific organisational decision.
- Vendor contracts do not address whether candidate data is used for model training.
- No internal complaints procedure exists or has been communicated to candidates.

Bias monitoring and vendor governance

Why this category matters

The ICO's 2024 audit made almost 300 recommendations to AI recruitment tool providers, with bias monitoring among the most significant concerns. Vendors sometimes define themselves incorrectly as processors rather than controllers, use vague contracts that obscure responsibilities, and in some cases use candidate data for model training without the recruiter's awareness. The EU AI Act requires deployers to monitor the performance of high-risk AI systems throughout their lifecycle, not only before launch. An organisation that audited its tool at deployment and has not reviewed it since is running on a governance foundation that may no longer hold.

Questions to review

1	CORE	Do we require pre-deployment bias testing from every AI hiring vendor, covering key protected characteristics?
2	CORE	Do our vendor contracts require notification when the model is updated materially?
3	CORE	Do our contracts include audit rights, explicit data processing instructions, and data deletion requirements?
4	CORE	Do we monitor hiring funnel outcomes by demographic group on an ongoing basis?
5	CORE	Is there a named internal person responsible for reviewing bias monitoring data?
6	DEEPER DIVE	Do we require vendors to provide model cards or equivalent technical documentation before deployment?
7	DEEPER DIVE	Are vendor performance reviews conducted annually, not only at contract renewal?

WHAT GOOD LOOKS LIKE

- Pre-deployment bias testing is a contractual requirement, covering race, gender, age, and disability at minimum.
- Vendor contracts require notification of model updates and trigger a re-evaluation when updates are made.
- Outcome monitoring by demographic group is conducted at least six-monthly and reviewed by a named person.

RED FLAGS AND WEAK PRACTICE

- Bias testing is accepted on vendor assurance with no independent evidence or contractual requirement.
- Vendors can update their model without notifying the organisation.
- No outcome monitoring takes place after deployment.
- Contracts contain no AI-specific provisions for audit rights, data instructions, or update notification.

Candidate rights and incident management

Why this category matters

The Data (Use and Access) Act 2025 updated UK candidate rights around automated decisions. Candidates have the right to be informed when AI significantly influences a decision, to contest that decision, and to request human review. A mandatory complaints procedure came into force on 19 June 2026. The EU AI Act requires deployers to report serious incidents involving high-risk AI to national competent authorities. Absent a clear escalation path, organisations respond to incidents reactively and inconsistently, increasing both legal and reputational exposure.

Questions to review

1	CORE	Are candidates told when AI is being used in their application, at which stage, and what it is assessing?
2	CORE	Is there a documented process for candidates to request human review of an AI-influenced decision?
3	CORE	Does our complaints procedure meet the DUAA 2025 requirement, including acknowledgement within 30 days?
4	CORE	Is there a clear escalation path from recruiter to HR Director to DPO to General Counsel for AI-related incidents?
5	CORE	Is there a named person with authority to pause or suspend an AI tool pending investigation?
6	DEEPER DIVE	Do we have trigger criteria that automatically initiate a review, such as bias metrics breaching thresholds or a candidate complaint?
7	DEEPER DIVE	For EU operations: is there a process for reporting serious AI incidents to national competent authorities?

WHAT GOOD LOOKS LIKE

- Candidates receive clear, specific disclosure before each AI-assessed stage, in plain language.
- A documented complaints procedure is operational, with acknowledgement within 30 days and a named contact.
- A tiered escalation path is documented and known to the people in each tier.
- A named person has the authority to pause a tool and that authority is documented.

RED FLAGS AND WEAK PRACTICE

- Candidates receive generic privacy policy wording with no meaningful information about AI involvement.
- No complaints procedure exists or has been communicated to candidates.
- Escalation is handled informally with no documented path or defined authority.
- Nobody has been named as having authority to pause an AI tool.

Policy readiness summary

Use this page to record your overall readiness across all six categories before completing the policy template. Identify your priority gaps and assign owners. Transfer these to the relevant sections of Part 2.

Category	Current status	Priority gap or action	Owner
Governance ownership and decision rights			
AI tool inventory and classification			
Human oversight standards			
Data protection obligations			
Bias monitoring and vendor governance			
Candidate rights and incident management			

Status guide: Complete (all CORE questions met) / Partial (some CORE gaps) / Gap (significant CORE gaps requiring immediate action)

AI Hiring Policy Template

This template is a starting point, not a finished policy. Every section includes sample text written in plain English and suitable for governance use. Replace the sample text with language that accurately reflects your organisation, your tools, and your actual process. Have the completed document reviewed by your legal team before final approval.

Policy document details

Organisation name	[Organisation name]
Document title	AI Hiring Policy
Version	[1.0]
Status	[Draft / Under review / Approved]
Policy owner	[Name, Role]
Approved by	[Name, Role]
Date approved	[Date]
Next review date	[Date]
Document owner	[Name, Role]
Date completed	[Date]

Purpose and scope

State clearly why this policy exists, what it covers, and what it does not cover. Keep it short. This section is the foundation everything else rests on.

1.1 Purpose

One paragraph. State what the policy is for and why it exists.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

[Organisation name] uses AI tools at certain stages of our hiring process. This policy sets out how we use those tools responsibly: who is accountable for them, what standards they must meet, how we protect candidates, and how we comply with our legal obligations. The goal is to ensure that AI supports our hiring decisions without replacing human judgement, and that every candidate is treated fairly and with transparency throughout the process.

1.2 Scope

List what the policy covers. Be specific about tools, stages, and geographies.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

This policy applies to all AI tools used in our hiring process, including: applicant tracking system features that filter, score, or rank candidates; AI-powered CV screening tools; video interview analysis platforms; automated assessment or psychometric tools; chatbots used for candidate pre-screening; and any generative AI tools used to support hiring decisions. It applies across all hiring in [list jurisdictions]. It applies to all employees involved in hiring, including HR, talent acquisition, hiring managers, and any third parties acting on our behalf.

1.3 Related policies

List the other policies this one connects to.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

This policy should be read alongside our Data Protection Policy, Equal Opportunities Policy, Recruitment and Selection Policy, and Supplier Code of Conduct.

Definitions

Define the key terms used in this policy. Plain language. Avoid legal drafting. The goal is that any recruiter or hiring manager can read these and understand what they mean in practice.

Key terms

Adapt these definitions to reflect your tools and context.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

AI system: software that uses machine learning, statistical, or rule-based techniques to produce outputs such as scores, rankings, classifications, or recommendations about candidates. Automated decision-making (ADM): a decision made with no meaningful human involvement, or significantly influenced by automated processing.

High-risk AI: under the EU AI Act, AI systems used for recruitment, selection, or employment management decisions are classified as high-risk, requiring a higher standard of governance and oversight. Meaningful human review: a human genuinely participates in the decision in a way that can influence the outcome. Viewing an AI score and accepting it without scrutiny does not constitute meaningful review. Significant decision: a decision that produces a legal or similarly significant effect for a candidate, such as rejection from a hiring process.

Governance and accountability

Name the people and roles responsible for this policy. This section must have real names or specific roles attached to it, not generic function labels.

3.1 Policy owner

Name the role accountable for maintaining and enforcing this policy.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

Policy owner: [CHRO / Head of Talent / HR Director] Executive sponsor: [CEO / COO / CPO] This policy owner is responsible for: approving new AI hiring tools; maintaining this policy and ensuring it is reviewed on schedule; escalating significant AI-related incidents to the executive sponsor.

3.2 Mandatory reviewers

List the roles who must be consulted when this policy is reviewed or when new AI tools are approved.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

The following roles must be consulted before any new AI hiring tool is approved for use, and must review this policy at each annual review: - DPO or Head of Data Protection - General Counsel or Head of Employment Law - CISO or Head of IT Security - DEI Lead - Head of Procurement

3.3 Decision rights

Specify who can approve, pause, or remove AI tools.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

Approving a new AI hiring tool: [Policy owner] with sign-off from DPO and General Counsel. Pausing an AI tool pending investigation: [Policy owner] or DPO. Removing an AI tool permanently: [Policy owner] with executive sponsor approval. Overriding an AI recommendation: Any trained recruiter or hiring manager. All overrides must be recorded in [name of system].

Approved and prohibited uses

Be specific about what AI can and cannot do in your hiring process. Vague policies create gaps that are exploited by default.

4.1 Approved uses

List the AI uses your organisation permits, with any conditions attached.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

The following uses of AI in hiring are approved, subject to the governance controls in this policy: - CV triage and initial ranking, provided a human reviews the shortlist before any rejection decision is made - Structured assessment platforms, provided bias testing has been completed and reasonable adjustments are available - AI-assisted interview scheduling and candidate communication - Drafting job descriptions and interview questions with generative AI tools, provided a human reviews and approves the output - Interview transcription tools, provided candidates are informed before the interview begins

4.2 Prohibited uses

List what is not permitted. These should be clear prohibitions, not guidelines.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

The following uses of AI in hiring are prohibited: - Fully automated rejection of candidates with no meaningful human review - AI tools designed to infer emotions, personality, or mental states from facial expression, voice, or behavioural signals - Inferring protected characteristics from application data, names, or photographs - Pasting candidate data into public generative AI tools not covered by a data processing agreement - Using any AI hiring tool not listed on our approved tool inventory

Human oversight

This is the section most likely to fail in practice. Make the required steps specific enough that a recruiter knows exactly what to do.

5.1 Meaningful human review

Define what review looks like in your process. Be specific.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

Before confirming any shortlisting or rejection decision that has been influenced by an AI tool, the reviewer must:

1. Read the candidate's underlying application or assessment, not only the AI score or summary
 2. Consider whether the AI output is consistent with the candidate's actual materials
 3. Record their decision in [name of system], noting whether they agreed with or departed from the AI recommendation
 4. Be prepared to explain their decision if asked by the candidate or a manager
- Viewing an AI score and accepting it without completing the steps above does not constitute meaningful review and is not compliant with this policy.

5.2 Candidate right to human review

State how candidates can request human review and what happens when they do.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

Any candidate who believes an AI-influenced decision has affected their application may request human review. Requests should be sent to [email address]. We will acknowledge the request within [number] working days and provide a substantive response within [number] working days. The reviewer will be a qualified member of the talent team who was not involved in the original AI-assisted assessment.

Data protection and candidate rights

Cover the key data protection requirements in plain language. This section is not a DPIA. It is a policy-level statement of your obligations and how you meet them.

6.1 Data protection requirements

State your core data protection commitments for AI in hiring.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

We will complete a Data Protection Impact Assessment (DPIA) before deploying any AI hiring tool. DPIAs will be reviewed when tools change materially. We will hold candidate data generated by AI tools only for as long as necessary: [specify retention periods for your tools]. We will not permit vendors to use candidate data to train their AI models unless we have a separate lawful basis and have explicitly agreed to this in the vendor contract.

6.2 Candidate rights and complaints

State the rights candidates have and how they exercise them.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

Candidates have the right to: - Be informed when AI is being used in their application and what it is assessing - Request human review of a significant decision influenced by automated processing - Challenge a decision they believe is inaccurate or unfair - Raise a complaint about how AI has been used in their application Complaints should be sent to [email address]. We will acknowledge all complaints within 30 days. Complaints will be handled by [named role], who will coordinate with our DPO where required.

Bias monitoring and vendor governance

State your minimum standards for vendor selection and ongoing monitoring. These should be specific enough to be included as requirements in vendor contracts.

7.1 Vendor requirements

State what vendors must provide before and after deployment.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

Before any AI hiring tool is approved for use, the vendor must provide: - Evidence of independent bias testing covering race, gender, age, and disability at minimum - A description of the training data used, including its demographic composition - A data processing agreement meeting ICO and, where applicable, EU AI Act requirements - Confirmation that candidate data will not be used to train their central AI models without a separate agreement Vendor contracts must include: notification of material model updates; audit rights; data deletion requirements at contract end; incident notification obligations.

7.2 Ongoing monitoring

State how you monitor AI tool performance after deployment.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

We will review the performance of each active AI hiring tool at least every [six months / twelve months]. Reviews will cover: bias and adverse impact across protected groups; candidate complaint data; any changes made by the vendor to the underlying model. Reviews will be led by [named role] and findings reported to [named role or governance group].

Incident management and escalation

Define what counts as an incident and what happens when one occurs. The escalation path must be specific enough to follow under pressure.

8.1 Escalation path

Name the tiers and the people or roles in each.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

Tier 1: Recruiter or HRBP handles routine candidate questions and minor anomalies. Tier 2: [Policy owner / Head of Talent] handles formal complaints, tool performance concerns, and decisions to pause a tool. Tier 3: DPO and General Counsel handle regulatory queries, significant bias findings, data incidents, and DPIA updates. Tier 4: Executive sponsor handles material regulatory investigations, significant reputational incidents, and tool discontinuation.

8.2 Trigger conditions

List the events that automatically require escalation or a policy review.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

The following events require immediate escalation to Tier 2 or above: - A candidate complaint alleging discrimination in an AI-assisted process - Bias monitoring data showing adverse impact on a protected group above agreed thresholds - A data incident involving candidate data held by a vendor - A vendor notifying us of a material model update - Media or regulatory enquiry about our use of AI in hiring - Any serious malfunction of an AI hiring tool affecting candidate outcomes

Training and awareness

State who must be trained and on what. Training is the mechanism that connects policy to practice.

Training requirements

Specify who needs training and what it must cover.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

All recruiters and hiring managers involved in AI-assisted stages of the hiring process must complete training before using any approved AI hiring tool. Training must cover: what the tool measures and how to interpret its output; what meaningful human review requires in practice; how to override an AI recommendation and record the rationale; how to respond to candidate questions about AI use; how to escalate concerns. Training must be refreshed annually and whenever a new AI tool is introduced. Completion is recorded in [HR system / L&D; platform].

Policy review and version control

State when and how this policy is reviewed. This section should be completed at the point of approval.

Review schedule

Set the review cycle and the trigger conditions for an early review.

Sample text

Adapt this to your organisation. Replace all placeholders in brackets.

This policy will be reviewed in full at least annually. An early review will be triggered by: any significant change to the EU AI Act or UK data protection framework; any material change to AI tools in use; any significant AI-related incident or regulatory enquiry; a recommendation from the DPO or General Counsel. Version: [1.0] Approved by: [Name, Role] Date approved: [Date] Next review date: [Date] Policy owner: [Name, Role]

Regulatory timeline

Key dates for UK and EU organisations deploying AI in hiring. Current as of June 2026. Review annually and update as new regulations come into force.

Date	Development	Jurisdiction
March 2024	DSIT Responsible AI in Recruitment guidance published	UK
March 2024	ICO recruitment ADM consultation closes	UK
1 August 2024	EU AI Act enters into force	EU/EEA
November 2024	ICO AI in Recruitment Outcomes Report published (296 recommendations)	UK
February 2025	ICO key data protection questions for AI in recruitment published	UK
2 February 2025	EU AI Act prohibitions apply, including prohibition on emotion inference at work	EU/EEA
5 February 2026	DUAA 2025 ADM provisions in force	UK
January 2026	CIPD HR role in AI governance guidance published	UK
April 2026	ICO AI recruitment fairness spotlight published	UK
19 June 2026	DUAA 2025 mandatory complaints procedure in force	UK
2 August 2026	EU AI Act Article 50 transparency rules apply (AI disclosure, content labelling)	EU/EEA
2 December 2027	EU AI Act high-risk AI requirements fully applicable, including recruitment and selection	EU/EEA

This timeline reflects the regulatory position as of June 2026. Requirements in this area are evolving. Review annually and seek legal advice on your specific obligations.

Final notes

This toolkit is a governance and policy development tool. It is designed to help organisations build a responsible AI hiring policy efficiently, starting from a structured review of their current position and finishing with a draft policy ready for legal review and board approval.

It is not legal advice. The policy template is a starting point, not a legally complete document. Have your completed policy reviewed by your legal team, DPO, and employment law counsel before final approval and publication.

AI hiring governance is no longer a future consideration. For many organisations it is becoming a current operational requirement. The purpose of this toolkit is not to slow adoption. It is to help organisations adopt AI in hiring responsibly, transparently, and with confidence.

Part of the AI Hiring Toolkit series

Toolkit 1	AI Procurement Questions for Hiring Teams	Live
Toolkit 2	Candidate Transparency Guide	Live
Toolkit 3	Bias Audit Checklist for AI Hiring	Live
Toolkit 4	AI Hiring Policy Framework	This document