

AI Procurement Questions for Hiring Teams

Choose AI hiring vendors with greater confidence, accountability, and transparency before signing a contract.

INCLUDES

- 48 structured evaluation questions
- Vendor comparison framework
- Internal alignment prompts
- Governance review sheet

AI Procurement Questions for Hiring Teams

What this toolkit is for	A structured set of vendor evaluation questions for HR and procurement teams assessing AI hiring tools. It gives non-technical buyers the language and framework to ask better questions, recognise weak answers, and make more informed procurement decisions.
Who it is for	HR Directors, Talent Acquisition leads, Procurement teams, Legal and Compliance stakeholders. Written for people who are not AI experts and should not need to be.
What it produces	A structured vendor evaluation across six categories, with a comparison sheet to record findings across multiple vendors. The output is a side-by-side record that supports a more defensible procurement decision.
Why it matters	AI hiring tools make consequential decisions about candidates. The vendor contract almost always protects the vendor more than the buyer. Most teams do not yet have a structure for asking the right questions before they sign. This toolkit closes that gap.
Position in the series	Toolkit 1 of 4 in the AI Hiring Toolkit series. The natural starting point: before you deploy a tool, before you build a transparency framework, before you audit for bias, and before you write a policy.

A vendor who cannot answer these questions clearly is answering you anyway.

AI Procurement Questions for Hiring Teams

How to use this toolkit

This toolkit is for HR Directors, Talent Acquisition leads, Procurement teams, and Legal or Compliance stakeholders evaluating AI hiring tools. It is written for people who are not AI experts. The goal is to give non-technical buyers the language and structure to ask better questions and recognise weak answers.

Step 1	Work through the internal alignment section with your team before any vendor conversation. These prompts surface assumptions and gaps that are better resolved before you are in a demo.
Step 2	Select the categories most relevant to your context. All six matter, but weight them based on your industry, candidate volume, and regulatory exposure.
Step 3	Use the questions live in vendor calls, not just as pre-sent documents. How a vendor responds in real time tells you as much as the answer itself.
Step 4	Use the comparison sheet after each vendor conversation to build a side-by-side record of your evaluation.

Question types

CORE	Prioritise in every vendor conversation. Ask these first.
DEEPER DIVE	For follow-up calls, or when Legal or Risk are in the room.

Estimated completion time	60 to 90 minutes per vendor, plus 30 minutes for team alignment before vendor conversations begin.
Primary owner	Head of Talent / HR Director / Procurement Lead
Suggested contributors	Legal or Employment Counsel, DPO or Data Protection Lead, DEI Lead, Hiring Manager or Business Owner
Suggested review cycle	Review when evaluating new vendors or renewing existing contracts. Update annually to reflect regulatory changes.

What this toolkit does not do

This is not a contract template, a legal opinion, or a substitute for legal review. It is a structured vendor evaluation tool. For contract negotiation, data processing agreements, and formal compliance review, work with qualified legal counsel.

Before you start

Internal alignment prompts

Do not skip this section. These prompts surface the gaps that matter most before you begin.

Work through these prompts with your team before the first vendor conversation. They are not counted in the 48 vendor questions. They exist to make sure your organisation arrives at procurement conversations with a clear internal position.

Do we have a named person accountable for AI use in our hiring process?

If the answer is unclear, vendor conversations will be harder to resolve. Someone needs to own what gets decided and what gets signed.

Do we understand which parts of our current process already involve AI, including tools embedded in our ATS?

Most organisations have more AI in their hiring process than they realise. Knowing your baseline makes vendor evaluation more honest.

Have we defined what we are trying to achieve with this tool, beyond saving time?

Faster screening is not a governance objective. Knowing what good looks like for your specific context makes it easier to evaluate vendor claims.

Do we know which protected characteristics are most relevant to our candidate population and hiring context?

Bias testing is only useful if it covers the characteristics that matter for your roles and candidate pool.

Have we discussed with Legal or Compliance what our obligations are under GDPR and AI regulation for this use case?

The EU AI Act classifies most AI hiring tools as high-risk systems. Your legal team should know what that means for procurement before you are in a contract.

Do we have a process for candidates to raise concerns about how AI is used in our hiring process?

If not, that is a gap to address regardless of which tool you buy. Some vendors will help you build this; ask them how.

Are our hiring managers and recruiters prepared to explain AI use to candidates if asked?

Candidates increasingly ask. A recruiter who cannot answer clearly creates a transparency problem even when the tool itself is sound.

Have we agreed internally on what meaningful human oversight looks like in our specific process?

Regulators and courts are clear that clicking approve without reviewing the underlying material is not meaningful oversight. Agree your standard before you deploy.

If your team cannot agree on the answers to these prompts, that is important information. It means you are not yet ready to evaluate vendors. Resolve internal alignment first.

Bias and fairness

Why this category matters

AI hiring tools trained on historical data can encode and scale existing biases in ways that are hard to detect after the fact. Most vendors will tell you they have tested for bias. The important question is how, with what methodology, across which protected characteristics, and what the results actually showed. This category helps you move past reassurance and into evidence.

Questions to ask

1	CORE	What protected characteristics does your bias testing cover, and how frequently is testing conducted?
2	CORE	Can you share the results of your most recent bias audit, including any disparate impact findings?
3	CORE	What methodology do you use to test for bias? Is it internal testing, third-party audit, or both?
4	CORE	If bias is identified in the outputs, what is your process for addressing it and notifying clients?
5	CORE	Has your tool ever been found to produce discriminatory outcomes in practice? If so, what happened?
6	DEEPER DIVE	How does the model perform across candidates with non-traditional backgrounds or career gaps?
7	DEEPER DIVE	What data was the model originally trained on, and how representative is it of diverse candidate populations?
8	DEEPER DIVE	Do clients receive ongoing fairness reporting, and how does performance vary across different role families?

WHAT GOOD LOOKS LIKE

- The vendor provides specific protected characteristics tested, shares a summary of results, names the methodology, and can describe what remediation looks like when issues are found.
- Results are available in writing, not just described verbally in a sales call.
- The vendor can explain what happens after bias is detected, including client notification.

RED FLAGS AND WEAK PRACTICE

- Vague assurances such as 'we take bias very seriously' without supporting evidence.
- Reluctance to share audit results or methodology.
- Claims that the tool is bias-free. This should raise concern rather than reassurance.
- Audits conducted only by the vendor on their own tool with no independent verification.

Data and privacy

Why this category matters

AI hiring tools process significant amounts of personal data about candidates, often including inferred attributes candidates are unaware of. Understanding what data is collected, how long it is retained, whether it trains future models, and who has access is not just a compliance question. It is a candidate trust question. This category surfaces practices most buyers do not ask about until something goes wrong.

Questions to ask

1	CORE	What personal data does your tool collect from candidates, and is any of it inferred rather than directly provided?
2	CORE	Is candidate data from our hiring process used to train or improve your models? Can we opt out?
3	CORE	How long is candidate data retained, and what is the deletion process when requested?
4	CORE	Which third parties or subprocessors have access to candidate data, including any external model providers?
5	CORE	How do you handle data subject access requests from candidates under GDPR or equivalent regulation?
6	DEEPER DIVE	Who within your organisation has access to candidate data from our account?
7	DEEPER DIVE	Where is candidate data stored geographically, and does it transfer outside the UK or EU?
8	DEEPER DIVE	What happens to our candidate data if we end our contract with you?

WHAT GOOD LOOKS LIKE

- Clear, specific answers. The vendor can point to their data processing agreement, has a documented deletion process, allows opt-out of training data use, and can name their subprocessors and data residency arrangements.
- Written confirmation that candidate data is not used for model training without explicit agreement.
- Specific retention timelines rather than references to industry-standard practices.

RED FLAGS AND WEAK PRACTICE

- Inability to confirm whether data is used for model training.
- Vague references to industry-standard security without specifics.
- No clear answer on data residency, deletion timelines, or which subprocessors have access.
- Data processing agreement is unavailable or not yet drafted.

Explainability and auditability

Why this category matters

If a candidate asks why they were not progressed, can you give them a meaningful answer? If a decision is challenged, can you reconstruct how it was reached? Many AI hiring tools produce outputs without a clear explanation of the inputs or logic involved. Explainability is both a data protection consideration and a basic standard of fairness. This category helps your legal and procurement teams assess the risks before you sign.

Questions to ask

1	CORE	How does your tool score or rank candidates, and what factors carry the most weight?
2	CORE	Which parts of the scoring process can you explain clearly, and which are less interpretable?
3	CORE	Can you explain a specific hiring decision in plain language if asked by a candidate or internal reviewer?
4	CORE	What documentation is produced for each candidate assessment, and how long is it retained?
5	DEEPER DIVE	What logs, decision records, or audit trails can we access directly as a client, rather than only on request?
6	DEEPER DIVE	How have you designed this tool to manage the risks around automated decision-making under data protection law?
7	DEEPER DIVE	If a candidate requests an explanation of a decision, what can you provide and in what timeframe?
8	DEEPER DIVE	Has your tool's decision-making process ever been subject to external audit? What did that involve?

WHAT GOOD LOOKS LIKE

- The vendor can describe in plain terms what the model scores and why. They have documentation practices in place, can support candidate explanation requests, and give clients direct access to audit trails.
- The vendor has considered their position on automated decision-making under data protection law and can articulate it.
- Candidate explanation requests have a defined process and timeline.

RED FLAGS AND WEAK PRACTICE

- Deflection on how the model works, citing proprietary IP.
- No process for candidate explanation requests.
- Uncertainty about data protection obligations.
- Claims that the model is just a tool and all decisions are human, while also describing automated shortlisting.

Human oversight and escalation

Why this category matters

Meaningful human oversight is not a checkbox that a recruiter clicks without reviewing the reasoning behind an automated output. This category establishes whether human review is genuinely embedded in the process, whether candidates have a route to escalate, whether recruiters can override the tool, and whether your organisation has defined accountability for what the system does in practice.

Questions to ask

1	CORE	At which specific points in the hiring process does human review take place?
2	CORE	Can a candidate request human review of an automated assessment? If so, how?
3	CORE	Can human reviewers override the tool's output, and is that override recorded for audit purposes?
4	CORE	Who in our organisation is accountable for decisions made with the assistance of this tool?
5	CORE	What safeguards exist to prevent hiring teams from rubber-stamping automated outputs?
6	DEEPER DIVE	If the tool's output contradicts a recruiter's assessment, what guidance do you provide?
7	DEEPER DIVE	What training do you provide or recommend for hiring teams using this tool?
8	DEEPER DIVE	How do you ensure human review points are genuinely substantive rather than procedural?

WHAT GOOD LOOKS LIKE

- The vendor can identify specific named points where human review occurs, has a candidate escalation process, and confirms overrides are tracked.
- Recruiter training is provided or recommended, covering how to interpret and challenge the tool's output.
- The vendor is clear that the tool supports rather than replaces human decision-making.

RED FLAGS AND WEAK PRACTICE

- Describing human in the loop without specifying where or how.
- No candidate escalation process.
- Framing the tool as autonomous in early stages.
- Inability to confirm whether overrides are tracked or what recruiter training looks like.

Vendor accountability

Why this category matters

When an AI hiring tool produces a bad outcome, who is responsible? Vendors often design contracts to limit their own liability while placing accountability on the buyer. This category helps you understand what the vendor will actually stand behind, what recourse exists when something goes wrong, and whether accountability is genuinely shared or quietly transferred to your organisation at the point of signing.

Questions to ask

1	CORE	How do you approach liability and client support if the tool contributes to a harmful or discriminatory outcome?
2	CORE	What is your incident response process if a bias, security, or model-performance issue is discovered?
3	CORE	How quickly would we be notified if an issue affecting our hiring process were identified?
4	CORE	What ongoing support do you provide to help clients use the tool responsibly and compliantly?
5	DEEPER DIVE	Do you carry professional indemnity or product liability insurance for AI-influenced hiring decisions?
6	DEEPER DIVE	Have any clients faced legal challenges related to the use of your tool? How were they handled?
7	DEEPER DIVE	What contractual protections exist if a tool failure leads to a decision that causes harm?
8	DEEPER DIVE	How do you track changes in employment law and AI regulation, and how do you communicate these to clients?

WHAT GOOD LOOKS LIKE

- Clear liability position, documented incident response with defined notification timelines, and relevant insurance coverage.
- The vendor can reference how real issues have been handled, not just hypothetical processes.
- Ongoing compliance support is available and documented.

RED FLAGS AND WEAK PRACTICE

- Immediate redirection to the legal team when liability is mentioned.
- Statements that all decisions are the client's responsibility.
- No incident response process or clear notification timeline when issues are discovered.

Contract and exit

Why this category matters

Buying an AI hiring tool is easier than leaving one. Data becomes embedded, processes are built around outputs, and exit terms are rarely scrutinised at signing. This category ensures your organisation understands what happens to your data if you leave, what portability looks like in practice, what transition support is available, and whether you are entering a relationship that is genuinely reversible.

Questions to ask

1	CORE	What happens to all candidate data held by your platform when we end the contract?
2	CORE	Can we export our historical data in a usable format before we leave? What does that look like?
3	CORE	What is the minimum contract term, and what are the notice requirements or penalties for early exit?
4	CORE	What support do you provide during offboarding or migration to another vendor?
5	DEEPER DIVE	Do you use our data or hiring outcomes to improve your models after we leave?
6	DEEPER DIVE	Are there technical or commercial arrangements that would make switching to a different vendor difficult?
7	DEEPER DIVE	What does the full off-boarding process look like, and how long does complete data deletion take?
8	DEEPER DIVE	If your company is acquired or ceases to operate, what protections exist for our data and continuity?

WHAT GOOD LOOKS LIKE

- A documented exit process with defined timelines for data deletion and export.
- Data available in portable formats. No punitive exit terms.
- A clear answer on post-contract data use. Transition support included or available.

RED FLAGS AND WEAK PRACTICE

- Vague deletion timelines. No exportable format for historical data.
- Long minimum terms with significant exit penalties.
- No transition support. Uncertainty about whether your data continues to train their models after you leave.

Vendor comparison sheet

Complete this sheet after each vendor conversation. Score each category 1 to 5, where 1 is poor or evasive and 5 is clear, specific, and evidenced.

Lead evaluator	[Name, Role]
Date completed	[Date]
Recommended vendor	[Vendor name, or 'None at this stage']
Key concerns	[Note any significant concerns that should inform the final decision]

On weighting: not all categories need equal weight. Adjust scoring emphasis to reflect your context. High-volume hiring may weight bias most. Regulated industries may weight explainability and contract terms most. Decide this before you compare vendors.

Category	Vendor A	Vendor B	Vendor C	Notes
Bias and fairness				
Data and privacy				
Explainability and auditability				
Human oversight and escalation				
Vendor accountability				
Contract and exit				
Overall score (out of 30)				
Recommended for next stage?				

Scoring guide: 5 = Clear, specific, evidenced. No follow-up needed. 4 = Good, minor gaps. 3 = Partial. Follow-up required. 2 = Vague or inconsistent. Significant concern. 1 = Evasive, unable to answer, or raised a red flag.

Final notes

This toolkit is a structured vendor evaluation tool. It is designed to make procurement conversations more productive and more honest, and to help organisations choose AI hiring tools with a clearer understanding of what they are buying.

It is not legal advice, a contract template, or a guarantee of compliance. For contract negotiation and data processing agreements, work with qualified legal counsel. The questions reflect current best practice in responsible AI procurement as of June 2026. Review and update your procurement process as the regulatory environment evolves.

Responsible procurement starts with better questions. The teams that ask them are already differentiated.

Part of the AI Hiring Toolkit series

Toolkit 1	AI Procurement Questions for Hiring Teams	This document
Toolkit 2	Candidate Transparency Guide	Live
Toolkit 3	Bias Audit Checklist for AI Hiring	Live
Toolkit 4	AI Hiring Policy Framework	Live