

Bias Audit Checklist

Identify where bias risk enters your hiring process and take practical steps to strengthen fairness and oversight.

INCLUDES

- 7 audit categories
- 40 review questions
- Remediation prompts
- Action planning framework

Bias Audit Checklist for AI Hiring

What this toolkit is for	A structured internal review framework for organisations using AI in their hiring process. It covers the seven categories where bias risk most commonly enters, scales, or goes undetected in AI-enabled hiring, with remediation prompts at each stage.
Who it is for	HR Directors, Heads of Talent, AI Governance leads, DPOs, and DEI leads. Anyone responsible for ensuring that AI tools in the hiring process are performing fairly and can be accounted for.
What it produces	A structured review across seven bias risk dimensions, with a summary audit sheet to record your position, identify priority gaps, and assign remediation actions. The output is an honest picture of where you are, so you can close the gaps that matter most.
Why it matters	Bias did not get removed from hiring when AI arrived. It got a new set of tools, a larger operational reach, and a harder-to-challenge wrapper. Most organisations have some pre-deployment audit in place. Far fewer have ongoing monitoring, a named governance owner, or contractual safeguards that hold vendors to account.
Position in the series	Toolkit 3 of 4. The natural third step: once you have evaluated a vendor and built your transparency framework, this toolkit helps you ensure the tool is performing fairly throughout its operational life.

Automation did not eliminate bias from hiring. It encoded historical patterns, replicated them at scale, and made them harder to challenge by giving them the appearance of objectivity. This checklist helps you identify where that is happening in your process.

Bias Audit Checklist for AI Hiring

How to use this checklist

This checklist is a structured internal review tool for organisations using AI in their hiring process. It covers seven audit dimensions: the areas where bias risk most commonly enters, scales, or goes undetected in AI-enabled hiring. Work through each dimension with the person or team responsible for AI governance in your organisation.

For each question, assess whether your organisation currently meets the standard described. Mark each as: Met, Partial, or Gap. Use the remediation prompts to plan your next actions. The goal is not a perfect score. It is an honest picture of where you are, so you can close the gaps that matter most.

Step 1	Read the 'Why this category matters' section before reviewing the questions.
Step 2	Work through the questions with whoever owns AI governance, procurement, and legal in your organisation.
Step 3	Mark each question: Met, Partial, or Gap. Be honest. A gap you name is one you can address.
Step 4	Prioritise the CORE questions. Every CORE gap is a minimum standard that needs closing.
Step 5	Use the remediation prompts to set specific actions and owners for each gap identified.

Question types

CORE	Minimum standard questions. Every CORE gap represents meaningful legal or reputational exposure and should be treated as a priority.
DEEPER DIVE	Extended audit questions. Important for organisations scaling AI use, operating in regulated markets, or preparing for external scrutiny.

Estimated completion time	90 to 120 minutes with your governance team. Allow additional time for remediation planning.
Primary owner	HR Director / Head of Talent / AI Governance Lead
Suggested contributors	DPO or Data Protection Lead, DEI Lead, Legal or Employment Counsel, Procurement, IT Security
Suggested review cycle	At least every six months for active high-risk AI tools. Full review whenever a vendor updates their model or a new tool is deployed.

What this checklist does not do

This checklist does not audit your AI tool directly. It does not replace an independent technical bias audit. It does not provide legal advice on GDPR, the EU AI Act, or employment discrimination law. It gives your team a structured framework for identifying where bias risk exists in your current process, and where to focus your remediation effort.

Before you start

Internal alignment prompts

Do not skip this section. These prompts surface the gaps that matter most before you begin.

Work through these prompts with your team before starting the seven categories. They are not counted in the question total.

Can we name the AI tool or tools we are using in our hiring process and describe what each one does?

If the answer involves uncertainty, that is itself a gap. You cannot audit what you cannot name.

Who in our organisation is responsible for what our AI hiring tools do in practice?

Not who procured them. Not who uses them day to day. Who owns the accountability for what they produce.

When were our AI hiring tools last reviewed for bias, and by whom?

If the answer is 'at deployment' or 'never', post-deployment monitoring is a priority dimension for your audit.

Do we know whether our vendor has updated their model since we last reviewed it?

Most vendor contracts do not require notification of model updates. Most organisations do not track them. This is the single most common governance gap in AI hiring.

If a candidate challenged an AI-influenced decision tomorrow, what would we do?

If the answer is uncertain, candidate impact and redress is a priority dimension.

Training data and model design

Why this category matters

Bias is most often introduced at the design stage, before a tool is ever deployed. A model trained on historical hiring data will encode the patterns of who was previously hired and promoted. If that population was not diverse, the model is not finding the best candidates. It is finding candidates who resemble the people previously hired into a system that had its own biases long before the AI arrived.

Questions to audit

1	CORE	Do we know what data the AI tool was trained on?
2	CORE	Has the vendor confirmed whether the training data reflected a historically diverse population?
3	CORE	Do we know which candidate attributes the model was designed to predict or assess?
4	CORE	Has the vendor explained how the model defines a successful candidate, and what historical data informed that definition?
5	DEEPER DIVE	Are the assessment criteria the model uses documented in writing?
6	DEEPER DIVE	Has anyone in our organisation reviewed whether those criteria could disadvantage protected groups?

WHAT GOOD LOOKS LIKE

- The vendor has provided a plain-language explanation of what the model was trained on and what it predicts.
- The training data has been reviewed for representation across protected characteristics.
- Assessment criteria are documented and have been reviewed for potential disparate impact.

RED FLAGS AND WEAK PRACTICE

- The vendor describes the model as proprietary and provides no training data information.
- We have never asked what the model was trained on.
- Assessment criteria are described in marketing terms rather than technical specifics.

REMEDIATION PROMPTS

- Request a written explanation of training data composition from your vendor.
- Ask the vendor directly: 'Was your training data reviewed for demographic representation?'
- Document the assessment criteria and review them against your own candidate population.

Pre-deployment bias testing

Why this category matters

A bias audit conducted before deployment is the minimum standard. It is not sufficient on its own, because models can drift after deployment and training data patterns may not represent your specific candidate pool. But an organisation that has not conducted any pre-deployment testing has no baseline from which to assess whether the tool is performing fairly.

Questions to audit

1	CORE	Has our AI hiring tool been independently audited for bias before or after deployment?
2	CORE	Do we know which protected characteristics were included in the audit?
3	CORE	Was the audit conducted by an independent party, or by the vendor testing their own tool?
4	CORE	Are the audit results available to us, at minimum in summary form?
5	DEEPER DIVE	Does the audit cover the candidate population we actually hire from, or a generic benchmark?
6	DEEPER DIVE	Has the audit been repeated since the tool was first deployed?

WHAT GOOD LOOKS LIKE

- An independent bias audit has been conducted and results are available.
- The audit covers race, gender, age, and disability at minimum.
- The audit methodology is documented and was not designed by the vendor.

RED FLAGS AND WEAK PRACTICE

- No bias audit has been conducted.
- Only a vendor-conducted self-assessment exists.
- An audit was conducted but the organisation has not reviewed the results.
- The audit covered generic populations that do not reflect your candidate base.

REMEDIATION PROMPTS

- Request the full bias audit report from your vendor, or a summary if full results are withheld.
- If no independent audit exists, commission one before further deployment at scale.
- Check whether your vendor's audit methodology meets EEOC or ICO guidance standards.

Protected characteristics coverage

Why this category matters

Not all bias audits are equal. An audit that tests for gender bias only is not an audit of your tool's fairness. The protected characteristics most commonly underrepresented in audits include age, disability, and neurodivergence. AI interview tools that assess speech, communication style, or behavioural signals are particularly likely to disadvantage neurodivergent candidates and those for whom English is a second language.

Questions to audit

1	CORE	Does our audit or testing cover racial and ethnic background?
2	CORE	Does it cover gender, including non-binary identities?
3	CORE	Does it cover age, including candidates over 50?
4	CORE	Does it cover disability, including non-visible disabilities?
5	CORE	Has anyone assessed whether the tool could disadvantage neurodivergent candidates?
6	DEEPER DIVE	If the tool assesses speech or communication style, has it been tested against candidates for whom English is a second language?
7	DEEPER DIVE	Does the audit cover intersectional effects, such as the interaction of race and gender?

WHAT GOOD LOOKS LIKE

- The audit covers at minimum: race, gender, age, and disability.
- Neurodivergent candidates have been specifically considered in tool design or testing.
- Any communication or speech-based tool has been tested across linguistic backgrounds.

RED FLAGS AND WEAK PRACTICE

- The audit covers only one or two protected characteristics.
- Neurodivergence has not been considered.
- The tool assesses communication style but has not been tested across neurological or linguistic diversity.

REMEDIATION PROMPTS

- Map your current audit coverage against the full list of protected characteristics.
- Raise neurodivergence explicitly with your vendor as a testing requirement.
- For speech or communication tools, request evidence of multilingual and neurodivergent testing.

Post-deployment monitoring

Why this category matters

A bias audit conducted at deployment does not remain valid indefinitely. Vendors update their models. Candidate pools change. Regulatory expectations shift. Gartner research from Q1 2026 found 64% of enterprises have no formal performance monitoring for their AI systems. NIST published a formal position paper on this gap in March 2026. An organisation that audited its tool at deployment and has not reviewed it since is running on a foundation that may no longer hold.

Questions to audit

1	CORE	Do we have a defined review schedule for our AI hiring tool, for example annually or every six months?
2	CORE	Does our vendor contract require disclosure when the model is updated?
3	CORE	Do we monitor outcomes across protected groups in our hiring funnel?
4	CORE	Is there a named person responsible for reviewing monitoring data?
5	DEEPER DIVE	Do we have a mechanism to detect if the tool's outcomes are changing over time?
6	DEEPER DIVE	Have we reviewed our tool's performance since it was first deployed?
7	DEEPER DIVE	Does our contract trigger a re-evaluation if the model is updated?

WHAT GOOD LOOKS LIKE

- A defined review schedule exists and is documented.
- The vendor contract requires notification of model updates.
- Outcomes are monitored by protected group on at least a quarterly basis.
- A named person owns the monitoring process.

RED FLAGS AND WEAK PRACTICE

- No review schedule exists.
- The vendor can update the model without notifying us.
- Outcomes are not monitored after deployment.
- Nobody has reviewed the tool's performance since it went live.

REMEDIATION PROMPTS

- Set a review date in the next quarter and assign a named owner.
- Review your vendor contract for model update disclosure clauses. Add them at the next renewal if absent.
- Begin tracking hiring funnel outcomes by protected group, even informally.

Accountability and governance

Why this category matters

The accountability gap in AI hiring is not a technology problem. It is a governance problem. Most organisations are running tools that make consequential decisions about people with no named owner, no review process, and no clear line of responsibility when something goes wrong. UK GDPR already requires documented human oversight for automated decisions, and that is in effect now. The EU AI Act's equivalent requirement, a named competent human oversight person for high-risk AI systems, applies from December 2027. Building toward it now, while there is time to do it properly, is still the safer position.

Questions to audit

1	CORE	Is there a named person in our organisation who is accountable for what our AI hiring tool does?
2	CORE	Does that person understand what the tool measures and how decisions are made?
3	CORE	Is there a documented escalation route if the tool produces outcomes that require human review?
4	CORE	Are candidates notified that AI is being used and given a route to raise concerns?
5	DEEPER DIVE	Does the person responsible for the tool have the authority to suspend its use if a problem is identified?
6	DEEPER DIVE	Is the governance structure documented, or does it exist only informally?
7	DEEPER DIVE	Does our governance cover what happens when the tool produces an outcome a candidate challenges?

WHAT GOOD LOOKS LIKE

- A named person or role is documented as the owner of AI governance in hiring.
- That person can explain what the tool does and has the authority to act on problems.
- A documented escalation and review process exists.

RED FLAGS AND WEAK PRACTICE

- Nobody has been named as responsible for the tool in practice.
- Accountability is shared between HR, legal, and procurement with no single owner.
- Governance exists only as a vague understanding of who would deal with a problem if one arose.

REMEDIATION PROMPTS

- Name a single person responsible for AI governance in hiring and document it.
- Brief that person on what the tool measures, what the audit status is, and what the escalation route is.
- Create a one-page governance summary: tool name, owner, audit date, review schedule, escalation contact.

Vendor transparency and contractual safeguards

Why this category matters

AI hiring vendors operate with a significant information advantage. They understand their models. They know which questions are hard to answer. A vendor who is evasive about training data, audit methodology, or model update processes is communicating something about the robustness of their product. Contractual safeguards are the mechanism that turns governance intentions into enforceable commitments.

Questions to audit

1	CORE	Does our vendor contract specify what data the tool uses and how it is processed?
2	CORE	Does the contract require the vendor to notify us of significant model updates?
3	CORE	Does the contract include provisions for bias audit and what happens if issues are found?
4	DEEPER DIVE	Are the vendor's audit results available to us contractually, not just on goodwill?
5	DEEPER DIVE	Does the contract specify what happens to candidate data after the process ends?
6	DEEPER DIVE	Are there exit terms that allow us to move to a different system if compliance issues arise?
7	DEEPER DIVE	Does the contract define the vendor's liability if the tool produces discriminatory outcomes?

WHAT GOOD LOOKS LIKE

- The contract includes model update disclosure, audit access, and data retention clauses.
- Exit terms are clear and do not lock the organisation in if compliance issues arise.
- Vendor liability for discriminatory outcomes is addressed, even if capped.

RED FLAGS AND WEAK PRACTICE

- The contract is a standard vendor agreement with no AI-specific provisions.
- Model updates can be made without notification.
- Audit access depends on vendor goodwill rather than contractual right.

REMEDIATION PROMPTS

- Review your current contract against these criteria at the next renewal.
- Request a contract addendum covering model update disclosure and audit access.
- Seek legal review of vendor liability clauses before the next renewal.

Candidate impact and redress

Why this category matters

The purpose of a bias audit is not only to satisfy regulators. It is to ensure that candidates are treated fairly. A framework that identifies bias risk but provides no mechanism for affected candidates to seek redress is incomplete. UK GDPR Article 22 gives individuals the right to challenge automated decisions. The EU AI Act requires a mechanism for individuals to raise concerns about high-risk AI systems. The Data (Use and Access) Act 2025 mandatory complaints procedure came into force on 19 June 2026.

Questions to audit

1	CORE	Are candidates informed that AI is being used before they participate in an AI-assessed stage?
2	CORE	Are candidates given a route to request human review of an AI-influenced decision?
3	CORE	Is that route communicated proactively, not only in response to a complaint?
4	CORE	Do we have a documented process for handling candidate challenges to AI decisions?
5	DEEPER DIVE	Can we tell a candidate, on request, what the AI assessed them on?
6	DEEPER DIVE	Is there a timeframe for responding to candidate challenges and is it documented?
7	DEEPER DIVE	Have we reviewed our candidate redress process against UK GDPR Article 22 requirements?

WHAT GOOD LOOKS LIKE

- Candidates are told AI is involved before participating.
- A named contact and process for human review requests is communicated to all candidates.
- Response timeframes for challenges are defined and documented.

RED FLAGS AND WEAK PRACTICE

- Candidates are not informed AI is being used.
- No human review request process exists or is communicated.
- Challenges are handled ad hoc with no documented process.

REMEDIATION PROMPTS

- Add AI disclosure to all candidate communications before AI-assessed stages.
- Document a human review request process and communicate it in all relevant touchpoints.
- Review UK GDPR Article 22 guidance and confirm your process meets its requirements.

Summary

Audit summary sheet

Use this page to record your overall assessment across all seven categories. Transfer your ratings from each section, note your priority gaps, and assign owners for remediation actions.

Owner	[Name, Role]
Date completed	[Date]
Review date	[Date]

Category	Met / Partial / Gap	Priority action	Owner
Training data and model design			
Pre-deployment bias testing			
Protected characteristics coverage			
Post-deployment monitoring			
Accountability and governance			
Vendor transparency and contractual safeguards			
Candidate impact and redress			

Weighting note: All seven categories are important, but dimensions 4 (post-deployment monitoring) and 5 (accountability and governance) are the most common sources of organisational exposure. Most organisations have some pre-deployment audit in place. Far fewer have ongoing monitoring or a named governance owner.

Top 3 priority actions

1.	Action:	Owner:	By when:
2.	Action:	Owner:	By when:
3.	Action:	Owner:	By when:

Final notes

This checklist is a governance and review tool. It is designed to help organisations identify bias risk in their AI-enabled hiring process and take practical steps to address it.

It is not a substitute for an independent technical bias audit, and it does not constitute legal advice on employment discrimination law, GDPR, or the EU AI Act. The legal environment for AI in hiring is moving quickly. This checklist reflects best practice as of June 2026. Review and update your audit process as new requirements come into force.

The purpose of this checklist is not to eliminate all bias risk. It is to help organisations identify where bias risk exists, understand where governance gaps remain, and take practical steps to improve over time.

Part of the AI Hiring Toolkit series

Toolkit 1	AI Procurement Questions for Hiring Teams	Live
Toolkit 2	Candidate Transparency Guide	Live
Toolkit 3	Bias Audit Checklist for AI Hiring	This document
Toolkit 4	AI Hiring Policy Framework	Live